

Game Derandomization

Samuel Epstein*

April 22, 2023

Abstract

The recently introduced method of derandomization provides bounds on the Kolmogorov complexity of solutions to problems such as k-sat and graph coloring. This is done by showing a simple randomized method produces a solution with positive probability. This overall method can be applied to games, where if a probabilistic agent beats the environment, then a simple deterministic agent can be shown to win as well. In this paper, we show multiple examples of game derandomization.

Introduction

We describe two simplified cybernetic agent models. For the first model, the agent \mathbf{p} and environment \mathbf{q} are defined as follows. The agent is a function $\mathbf{p} : (\mathbb{N} \times \mathbb{N})^* \rightarrow \mathbb{N}$, where if $\mathbf{p}(w) = a$, $w \in (\mathbb{N} \times \mathbb{N})^*$ is a list of the previous actions of the agent and the environment, and $a \in \mathbb{N}$ is the action to be performed. The environment is of the form $\mathbf{q} : (\mathbb{N} \times \mathbb{N})^* \times \mathbb{N} \rightarrow \mathbb{N} \cup \{\mathbf{W}\}$, where if $\mathbf{q}(w, a) = b \in \mathbb{N}$, then b is \mathbf{q} 's response to the agent's action a , given history w , and the game continues. If \mathbf{q} responds \mathbf{W} then the agent wins and the game halts. The agent can be randomized. The game can continue forever, given certain agents and environments. This is called a win/no-halt game.

Theorem 1 ([Eps22]) *If probabilistic agent \mathbf{p}' wins against environment \mathbf{q} with at least probability p , then there is a deterministic agent \mathbf{p} of complexity $<^{\log} \mathbf{K}(\mathbf{p}') - \log p + \mathbf{I}(\langle p, \mathbf{p}', \mathbf{q} \rangle; \mathcal{H})$ that wins against \mathbf{q} .*

The mutual information term is $\mathbf{I}(x; \mathcal{H}) = \mathbf{K}(x) - \mathbf{K}(x|\mathcal{H})$, where \mathbf{K} is the prefix Kolmogorov complexity, $\mathcal{H} \in \{0, 1\}^\infty$ is the halting sequence, and for nonnegative function f , $<^{\log} f$ is defined to be $< f + O(\log(f+1))$. The second game is modified such that the environment gives a nonnegative rational penalty term to the agent at each round. Furthermore, the environment specifies an end to the game without specifying a winner or loser. This is called a penalty game.

Corollary 1 ([Eps22]) *If given probabilistic agent \mathbf{p} , environment \mathbf{q} halts with probability 1, and \mathbf{p} has expected penalty less than $n \in \mathbb{N}$, then there is a deterministic agent of complexity $<^{\log} \mathbf{K}(\mathbf{p}) + \mathbf{I}(\langle \mathbf{p}, n, \mathbf{q} \rangle; \mathcal{H})$ that receives penalty $< 2n$ against \mathbf{q} .*

*JP Theory Group. samepst@jpththeorygroup.org

Contents

1	EVEN-ODDS	2
2	GRAPH-NAVIGATION	2
3	PENALTY-TESTS	3
4	SET-SUBSET	4
5	MIN-CUT	4
6	COVER-TIME	5

1 EVEN-ODDS

We define the following win/no-halt game, entitled EVEN-ODDS. There are N rounds. At round 1, the environment \mathbf{q} secretly records bit $e_1 \in \{0, 1\}$. It sends an empty message to the agent who responds with bit $a_1 \in \{0, 1\}$. The agent gets a point if $e_1 \oplus a_1 = 1$. Otherwise the agent loses a point. For round i , the environment secretly selects a bit e_i that is a function of the previous agent's actions $\{a_j\}_{j=1}^{i-1}$ and sends an empty message to the agent, which responds with a_i and the agent gets a point if $e_i \oplus a_i = 1$, otherwise it loses a point. The agent wins after N rounds if it has a score of at least \sqrt{N} .

Theorem 2 *For large enough N , there is a deterministic agent \mathbf{p} that can win EVEN-ODDS with N rounds, with complexity $\mathbf{K}(\mathbf{p}) <^{\log} \mathbf{I}(\mathbf{q}; \mathcal{H})$.*

Proof. We describe a probabilistic agent \mathbf{p}' . At round i , \mathbf{p}' submits 0 with probability $1/2$. Otherwise it submits 1. By the central limit theorem, for large enough N , the score of the probabilistic agent divided by \sqrt{N} is $S \sim \mathcal{N}(0, 1)$. Let $\Phi(x) = \Pr[S > x]$. A common bound for $\Phi(x)$ is

$$\begin{aligned}\Phi(x) &> \frac{1}{2\pi} \frac{x}{x^2 + 1} e^{-x^2/2} \\ \Phi(1) &> \frac{1}{4\pi} e^{-1/2} > \frac{1}{8\pi}.\end{aligned}$$

Thus when $S \geq 1$, the score is at least \sqrt{N} . Thus \mathbf{p}' wins with probability at least $p = \frac{1}{8\pi}$. Thus by Theorem 1, there exists a deterministic agent \mathbf{p} that can beat \mathbf{q} with complexity

$$\mathbf{K}(\mathbf{p}) <^{\log} \mathbf{K}(\mathbf{p}') - \log p + \mathbf{I}((p, \mathbf{p}', \mathbf{q}); \mathcal{H}) <^{\log} \mathbf{I}(\mathbf{q}; \mathcal{H}).$$

□

2 GRAPH-NAVIGATION

The win/no-halt game is as follows. The environment \mathbf{q} consists of (G, s, r) . $G = (E, V)$ is a non-bipartite graph with undirected edges, $s \in V$ is the starting vertex, and $r \in V$ is the goal vertex. Let t_G be the time it takes for any random walk starting anywhere to converge to the stationary distribution $\pi(v)$, for all $v \in V$, up to a factor of 2.

There are t_G rounds and the agent starts at $s \in V$. At round 1, the environment gives the agent the degree $s \in V$, $\text{Deg}(s)$. The agent picks a number between 1 and $\text{Deg}(s)$ and sends it to \mathbf{q} . The agent moves along the edge the number is mapped to and is given the degree of the next vertex it is on. Each round's mapping of numbers to edges to be a function of the agent's past actions. This process is repeated t_G times. The agent wins if it is on $r \in V$ at the end of round t_G .

Theorem 3 *There is a deterministic agent \mathbf{p} that can win the GRAPH-NAVIGATION game with complexity $\mathbf{K}(\mathbf{p}) <^{\log} \log |E| + \mathbf{I}((G, s, r); \mathcal{H})$.*

Proof. It is well known if G is non-bipartite, a random walk starting from any vertex will converge to a stationary distribution $\pi(v) = \text{deg}(v)/2|E|$, for each $v \in V$.

A probabilistic agent \mathbf{p}' is defined as selecting each edge with equal probability. After t_G rounds, the probability that \mathbf{p}' is on the goal r is close to the stationary distribution π . More specifically the probability is $>^* 1/|E|$. Thus by Theorem 1, there is a deterministic agent \mathbf{p} that can find r in t_G turns and has complexity $\mathbf{K}(\mathbf{p}') <^{\log} \log |E| + \mathbf{I}((G, s, t); \mathcal{H})$.

3 PENALTY-TESTS

An example penalty game is as follows. The environment \mathbf{q} plays a game for N rounds, for some very large $N \in \mathbb{N}$, with each round starting with an action by \mathbf{q} . At round i , the environment gives, to the agent, an encoding of a program to compute a probability P_i over \mathbb{N} . The choice of P_i can be a computable function of i and the agent's previous turns. The agent responds with a number $a_i \in \mathbb{N}$. The environment gives the agent a penalty of size $T_i(a_i)$, where $T_i : \mathbb{N} \rightarrow \mathbb{Q}_{\geq 0}$ is a computable test, with $\sum_{a \in \mathbb{N}} P_i(a)T_i(a) < 1$. After N rounds, \mathbf{q} halts.

Theorem 4 *There is a deterministic agent \mathbf{p} that can receive a penalty $< 2N$ and has complexity $\mathbf{K}(\mathbf{p}) <^{\log} \mathbf{I}(\mathbf{q}; \mathcal{H})$.*

Proof. A very successful probabilistic agent \mathbf{p}' can be defined. Its algorithm is simple. On receipt of a program to compute P_i , the agent randomly samples a number \mathbb{N} according to P_i . At each round the expected penalty is $\sum_a P_i(a)T_i(a) < 1$, so the expected penalty of \mathbf{p} for the entire game is $< N$. Thus by Theorem 1, there is a deterministic agent \mathbf{p} such that

1. The agent \mathbf{p} receives a penalty of $< 2N$,
2. $\mathbf{K}(\mathbf{p}) <^{\log} \mathbf{I}(\mathbf{q}; \mathcal{H})$.

□

Let \mathbf{q} be defined so that $P_i(a) = [a \leq 2^i]2^{-i}$ and $T_i = [a \leq 2^i]2^{i-\mathbf{K}(a|i)}$, where $[A] = 1$ if A is true, and 0 otherwise. Thus each T_i is a randomness deficiency function. The probabilistic algorithm \mathbf{p}' will receive an expected penalty $< N$. However any deterministic agent \mathbf{p} that receives a penalty $< 2N$ must be very complex, as it must select many numbers with low randomness deficiency. Thus, by the bounds above, $\mathbf{I}(\mathbf{q}; \mathcal{H})$ must be very high. This makes sense because \mathbf{q} encodes N randomness deficiency functions.




Figure 1: A graphical depiction of a minimum cut. By removing the edges along the dotted line, two components are created.

4 SET-SUBSET

We define the following win/no-halt game, entitled SET-SUBSET. There are k rounds. At round $i = \{1, \dots, k\}$, the environment \mathbf{q} gives n numbers $A_i \subset \mathbb{N}$ to the agent \mathbf{p} . The environment secretly selects $m \leq n$ numbers $B_i \subseteq A_i$. The player selects a number $a_i \in A_i$. Each A_i and B_i are a function of the player's previous actions. The player wins if for every round, his selection a_i is in the secret set B_i . So for all $i \in \{1, \dots, k\}$, $a_i \in B_i$.

Theorem 5 *There is a deterministic agent \mathbf{p} that win against SET-SUBSET environment \mathbf{q} , where $\mathbf{K}(\mathbf{p}) <^{\log} k \log(n/m) + \mathbf{I}(\mathbf{q}; \mathcal{H})$.*

Proof. Let \mathbf{p}' be the randomized the player that selects a member of the given set with A_i with uniform probability. The probability that \mathbf{p}' picks a member of B_i is $|B_i|/|A_i| = m/n$. The probability that \mathbf{p}' picks a member of B_i for all $i \in \{1, \dots, k\}$ is $(m/n)^k$, which is the probability that \mathbf{p}' wins. $\mathbf{K}(\mathbf{p}') = O(1)$. Thus by Theorem 1, there exists a deterministic player \mathbf{p} that wins against \mathbf{q} with complexity bounded by the theorem statement. \square

5 MIN-CUT

We define the following win/no-halt game, entitled MIN-CUT. The game is defined by an undirected graph G and a mapping ℓ from numbers to edges. At round i , the environment \mathbf{q} sends the number of edges of G . The player responds with a number. The environment maps the number to an edge, and this mapping can be dependent on the player's previous actions. The environment then contracts the graph G along the edge. The game halts when the graph G has contracted into two vertices. The player wins if the cut represented by the contractions is a min cut. A minimum cut of a graph is the minimum number of edges, that when removed from the graph, produces two components. A graphical depiction of a min cut can be seen in Figure 1.

Theorem 6 *There is a deterministic agent \mathbf{p} that can win against MIN-CUT instance (G, S, ℓ) , $|G| = n$, such that $\mathbf{K}(\mathbf{p}) <^{\log} 2 \log n + \mathbf{I}((G, \ell); \mathcal{H})$.*

Proof. We define the following randomized agent \mathbf{p}' . At each round, \mathbf{p}' chooses an edge at random. Thus the interactions of \mathbf{p}' and \mathbf{q} represent an implementation of Karger's algorithm.

Karger’s algorithm has an $\Omega(1/n^2)$ probability of returning a min-cut. Thus \mathbf{p}' has an $\Omega(1/n^2)$ chance of winning. By Theorem 1, there exist a deterministic agent \mathbf{p} and c where \mathbf{p} can beat \mathbf{q} and has complexity $\mathbf{K}(\mathbf{p}) <^{\log} \mathbf{K}(\mathbf{p}') - \log c/n^2 + \mathbf{I}(\mathbf{q}; \mathcal{H}) <^{\log} 2 \log n + \mathbf{I}((G, \ell); \mathcal{H})$. \square

6 COVER-TIME

We define the following interactive penalty game. Let $G = (E, V)$ be a graph consisting of n vertices V and undirected edges E . The environment \mathbf{q} consists of (G, s, ℓ) . $G = (E, V)$ is a non-bipartite graph with undirected edges, $s \in V$ is the starting vertex. ℓ is a mapping from numbers to edges to be described later.

The agent starts at $s \in V$. At round 1, the environment gives the agent the degree $s \in V$, $\text{Deg}(s)$. The agent picks a number between 1 and $\text{Deg}(s)$ and sends it to \mathbf{q} . The agent moves along the edge the number is mapped to and is given the degree of the next vertex it is on. Each round’s mapping of numbers to edges, ℓ , is a computable function of the agent’s past actions. The game stops if the agent has visited all vertices and the penalty is the number of turns the agents takes.

Theorem 7 *There is a deterministic agent \mathbf{p} that can play against COVER-TIME instance (G, S, ℓ) , $|G| = n$, and achieve penalty $\frac{8}{27}n^3 + o(n^3)$ and $\mathbf{K}(\mathbf{p}) <^{\log} \mathbf{I}((G, s, \ell); \mathcal{H})$.*

Proof. A probabilistic agent \mathbf{p}' is defined as selecting each edge with equal probability. Thus the agent performs a random walk. The game halts with probability 1. Due to [Fei95], the expected time (i.e. expected penalty) it takes to reach all vertices is $\frac{4}{27}n^3 + o(n^3)$. Thus by Theorem 1 there is a deterministic agent \mathbf{p} that can reach each vertex with a penalty of $\frac{8}{27}n^3 + o(n^3)$ and has complexity

$$\mathbf{K}(\mathbf{p}) <^{\log} \mathbf{K}(\mathbf{p}') + \mathbf{I}((G, s, \ell); \mathcal{H}) <^{\log} \mathbf{I}((G, s, \ell); \mathcal{H}).$$

\square

References

- [Eps22] S. Epstein. The outlier theorem revisited. *CoRR*, abs/2203.08733, 2022.
- [Fei95] U Feige. A tight upper bound on the cover time for random walks on graphs. *Random Struct. Algorithms*, 6(1):51–54, 1995.